

Yelp Review Dataset Analysis

Spark Setup

In [1]:

```
%info
```

```
Current session configs: {'conf': {'spark.pyspark.python': 'python3',
'spark.pyspark.virtualenv.enabled': 'true',
'spark.pyspark.virtualenv.type': 'native',
'spark.pyspark.virtualenv.bin.path': '/usr/bin/virtualenv'}, 'kind':
'pyspark'}
```

ID	YARN Application ID	Kind	State	Spark UI	Driver log	Current session?
6	application_1619310664737_0007	pyspark	idle	Link	Link	

In [2]:

```
sc.list_packages()
```

Starting Spark application

ID	YARN Application ID	Kind	State	Spark UI	Driver log	Current session?
7	application_1619310664737_0008	pyspark	idle	Link	Link	✓

SparkSession available as 'spark'.

Package	Version
beautifulsoup4	4.8.1
boto	2.49.0
jmespath	0.9.4
lxml	4.4.2
mysqlclient	1.4.6
nltk	3.4.5
nose	1.3.4
numpy	1.14.5
pip	21.1
py-dateutil	2.2
python36-sagemaker-pyspark	1.2.6
pytz	2019.3
PyYAML	3.11
setuptools	56.0.0
six	1.13.0
soupsieve	1.9.5
wheel	0.36.2
windmill	1.6

In [3]:

```
sc.install_pypi_package("pandas==1.0.3")
```

Collecting pandas==1.0.3

Using cached pandas-1.0.3-cp36-cp36m-manylinux1_x86_64.whl (10.0 MB)

Collecting python-dateutil>=2.6.1

Using cached python_dateutil-2.8.1-py2.py3-none-any.whl (227 kB)

```
Requirement already satisfied: numpy>=1.13.3 in /usr/local/lib64/python3.6/site-packages (from pandas==1.0.3) (1.14.5)
Requirement already satisfied: pytz>=2017.2 in /usr/local/lib/python3.6/site-packages (from pandas==1.0.3) (2019.3)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.6/site-packages (from python-dateutil>=2.6.1->pandas==1.0.3) (1.13.0)
Installing collected packages: python-dateutil, pandas
Successfully installed pandas-1.0.3 python-dateutil-2.8.1
```

In [4]:

```
sc.install_pypi_package('matplotlib==3.2.1')
```

```
Collecting matplotlib==3.2.1
  Using cached matplotlib-3.2.1-cp36-cp36m-manylinux1_x86_64.whl (12.4 MB)
Collecting pyparsing!=2.0.4,!=2.1.2,!=2.1.6,>=2.0.1
  Using cached pyparsing-2.4.7-py2.py3-none-any.whl (67 kB)
Requirement already satisfied: numpy>=1.11 in /usr/local/lib64/python3.6/site-packages (from matplotlib==3.2.1) (1.14.5)
Requirement already satisfied: python-dateutil>=2.1 in /mnt/tmp/1619392843364-0/lib/python3.6/site-packages (from matplotlib==3.2.1) (2.8.1)
Collecting kiwisolver>=1.0.1
  Using cached kiwisolver-1.3.1-cp36-cp36m-manylinux1_x86_64.whl (1.1 MB)
Collecting cyclor>=0.10
  Using cached cyclor-0.10.0-py2.py3-none-any.whl (6.5 kB)
Requirement already satisfied: six in /usr/local/lib/python3.6/site-packages (from cyclor>=0.10->matplotlib==3.2.1) (1.13.0)
Installing collected packages: pyparsing, kiwisolver, cyclor, matplotlib
Successfully installed cyclor-0.10.0 kiwisolver-1.3.1 matplotlib-3.2.1 pyparsing-2.4.7
```

In [5]:

```
sc.install_pypi_package('seaborn==0.11.1')
```

```
Collecting seaborn==0.11.1
  Using cached seaborn-0.11.1-py3-none-any.whl (285 kB)
Requirement already satisfied: pandas>=0.23 in /mnt/tmp/1619392843364-0/lib64/python3.6/site-packages (from seaborn==0.11.1) (1.0.3)
Collecting scipy>=1.0
  Using cached scipy-1.5.4-cp36-cp36m-manylinux1_x86_64.whl (25.9 MB)
Requirement already satisfied: matplotlib>=2.2 in /mnt/tmp/1619392843364-0/lib64/python3.6/site-packages (from seaborn==0.11.1) (3.2.1)
Collecting numpy>=1.15
  Using cached numpy-1.19.5-cp36-cp36m-manylinux2010_x86_64.whl (14.8 MB)
Requirement already satisfied: pyparsing!=2.0.4,!=2.1.2,!=2.1.6,>=2.0.1 in /mnt/tmp/1619392843364-0/lib/python3.6/site-packages (from matplotlib>=2.2->seaborn==0.11.1) (2.4.7)
Requirement already satisfied: python-dateutil>=2.1 in /mnt/tmp/1619392843364-0/lib/python3.6/site-packages (from matplotlib>=2.2->seaborn==0.11.1) (2.8.1)
Requirement already satisfied: kiwisolver>=1.0.1 in /mnt/tmp/1619392843364-0/lib64/python3.6/site-packages (from matplotlib>=2.2->seaborn==0.11.1) (1.3.1)
Requirement already satisfied: cyclor>=0.10 in /mnt/tmp/1619392843364-0/lib/python3.6/site-packages (from matplotlib>=2.2->seaborn==0.11.1) (0.10.0)
Requirement already satisfied: six in /usr/local/lib/python3.6/site-packages (from cyclor>=0.10->matplotlib>=2.2->seaborn==0.11.1) (1.13.0)
Requirement already satisfied: pytz>=2017.2 in /usr/local/lib/python3.6/site-packages (from pandas>=0.23->seaborn==0.11.1) (2019.3)
Installing collected packages: numpy, scipy, seaborn
  Attempting uninstall: numpy
    Found existing installation: numpy 1.14.5
    Not uninstalling numpy at /usr/local/lib64/python3.6/site-packages, outside environment /mnt/tmp/1619392843364-0
    Can't uninstall 'numpy'. No files were found to uninstall.
```

Successfully installed numpy-1.19.5 scipy-1.5.4 seaborn-0.11.1

ERROR: pip's dependency resolver does not currently take into account all the packages that are installed. This behaviour is the source of the following dependency conflicts.

python36-sagemaker-pyspark 1.2.6 requires pyspark==2.3.2, which is not installed.

In [6]:

```
sc.list_packages()
```

Package	Version
beautifulsoup4	4.8.1
boto	2.49.0
cycler	0.10.0
jmespath	0.9.4
kiwisolver	1.3.1
lxml	4.4.2
matplotlib	3.2.1
mysqlclient	1.4.6
nltk	3.4.5
nose	1.3.4
numpy	1.19.5
pandas	1.0.3
pip	21.1
py-dateutil	2.2
pyparsing	2.4.7
python-dateutil	2.8.1
python36-sagemaker-pyspark	1.2.6
pytz	2019.3
PyYAML	3.11
scipy	1.5.4
seaborn	0.11.1
setuptools	56.0.0
six	1.13.0
soupsieve	1.9.5
wheel	0.36.2
windmill	1.6

Imports

In [7]:

```
from pyspark.sql.functions import explode, split

import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from pyspark.sql.functions import countDistinct
from pyspark.sql.functions import col, mean
```

Read

In [8]:

```
business = spark.read.json('s3://sta9760yelp/yelp_academic_dataset_business.json')
review = spark.read.json('s3://sta9760yelp/yelp_academic_dataset_review.json')
user = spark.read.json('s3://sta9760yelp/yelp_academic_dataset_user.json')
```

Overview of Data

```
In [9]: business.printSchema()
```

```
root
|-- address: string (nullable = true)
|-- attributes: struct (nullable = true)
|   |-- AcceptsInsurance: string (nullable = true)
|   |-- AgesAllowed: string (nullable = true)
|   |-- Alcohol: string (nullable = true)
|   |-- Ambience: string (nullable = true)
|   |-- BYOB: string (nullable = true)
|   |-- BYOBCorkage: string (nullable = true)
|   |-- BestNights: string (nullable = true)
|   |-- BikeParking: string (nullable = true)
|   |-- BusinessAcceptsBitcoin: string (nullable = true)
|   |-- BusinessAcceptsCreditCards: string (nullable = true)
|   |-- BusinessParking: string (nullable = true)
|   |-- ByAppointmentOnly: string (nullable = true)
|   |-- Caters: string (nullable = true)
|   |-- CoatCheck: string (nullable = true)
|   |-- Corkage: string (nullable = true)
|   |-- DietaryRestrictions: string (nullable = true)
|   |-- DogsAllowed: string (nullable = true)
|   |-- DriveThru: string (nullable = true)
|   |-- GoodForDancing: string (nullable = true)
|   |-- GoodForKids: string (nullable = true)
|   |-- GoodForMeal: string (nullable = true)
|   |-- HairSpecializesIn: string (nullable = true)
|   |-- HappyHour: string (nullable = true)
|   |-- HasTV: string (nullable = true)
|   |-- Music: string (nullable = true)
|   |-- NoiseLevel: string (nullable = true)
|   |-- Open24Hours: string (nullable = true)
|   |-- OutdoorSeating: string (nullable = true)
|   |-- RestaurantsAttire: string (nullable = true)
|   |-- RestaurantsCounterService: string (nullable = true)
|   |-- RestaurantsDelivery: string (nullable = true)
|   |-- RestaurantsGoodForGroups: string (nullable = true)
|   |-- RestaurantsPriceRange2: string (nullable = true)
|   |-- RestaurantsReservations: string (nullable = true)
|   |-- RestaurantsTableService: string (nullable = true)
|   |-- RestaurantsTakeOut: string (nullable = true)
|   |-- Smoking: string (nullable = true)
|   |-- WheelchairAccessible: string (nullable = true)
|   |-- WiFi: string (nullable = true)
|-- business_id: string (nullable = true)
|-- categories: string (nullable = true)
|-- city: string (nullable = true)
|-- hours: struct (nullable = true)
|   |-- Friday: string (nullable = true)
|   |-- Monday: string (nullable = true)
|   |-- Saturday: string (nullable = true)
|   |-- Sunday: string (nullable = true)
|   |-- Thursday: string (nullable = true)
|   |-- Tuesday: string (nullable = true)
|   |-- Wednesday: string (nullable = true)
|-- is_open: long (nullable = true)
```

```
-- latitude: double (nullable = true)
-- longitude: double (nullable = true)
-- name: string (nullable = true)
-- postal_code: string (nullable = true)
-- review_count: long (nullable = true)
-- stars: double (nullable = true)
-- state: string (nullable = true)
```

In [10]:

```
business.select('business_id','name','city','state','stars','categories').show(5)
business.count()
```

```
+-----+-----+-----+-----+-----+
| business_id | name | city | state | stars | categories |
+-----+-----+-----+-----+-----+
| 6iYb2HFDywm3zjuRg... | Oskar Blues Taproom | Boulder | CO | 4.0 | Gastropubs, Food, ... |
| tCbdrRPZA0oiIYSmH... | Flying Elephants ... | Portland | OR | 4.0 | Salad, Soup, Sand... |
| bvN78f1M8NLprQ1a1... | The Reclaimory | Portland | OR | 4.5 | Antiques, Fashion... |
| oaepsyvc0Jl7qwi8c... | Great Clips | Orange City | FL | 3.0 | Beauty & Spas, Hair... |
| PE9uqAjdW0E4-8mjG... | Crossfit Terminus | Atlanta | GA | 4.0 | Gyms, Active Life... |
+-----+-----+-----+-----+-----+
```

only showing top 5 rows

160585

Analyzing Categories

The categories column shows a list of categories the business falls under. We need to break out these categories to further analyze.

original

business_id	categories
RestaurantA	Spicy, Soup, Healthy

new

business_id	categories
RestaurantA	Spicy
RestaurantA	Soup
RestaurantA	Healthy

In [11]:

```
#creating new business df with categories broken out
```

```
business_cat = business.select('business_id','categories')
business_cat_explode = business_cat.withColumn('categories',explode(split('catego
```

```
In [12]: business_cat_explode.show(15)
business_cat_explode.count()
```

business_id	categories
6iYb2HFDywm3zjuRg...	Gastropubs
6iYb2HFDywm3zjuRg...	Food
6iYb2HFDywm3zjuRg...	Beer Gardens
6iYb2HFDywm3zjuRg...	Restaurants
6iYb2HFDywm3zjuRg...	Bars
6iYb2HFDywm3zjuRg...	American (Tradit...
6iYb2HFDywm3zjuRg...	Beer Bar
6iYb2HFDywm3zjuRg...	Nightlife
6iYb2HFDywm3zjuRg...	Breweries
tCbdrRPZA0oiIYSmH...	Salad
tCbdrRPZA0oiIYSmH...	Soup
tCbdrRPZA0oiIYSmH...	Sandwiches
tCbdrRPZA0oiIYSmH...	Delis
tCbdrRPZA0oiIYSmH...	Restaurants
tCbdrRPZA0oiIYSmH...	Cafes

only showing top 15 rows

708968

the number of rows increased as expected after breaking out categories

Total Unique Categories

```
In [13]: business_cat_explode.select(countDistinct('categories')).show()
```

count(DISTINCT categories)
2487

Top Categories By Business

Purpose is to find the most popular business categories. This will be done by grouping and finding the count of each category.

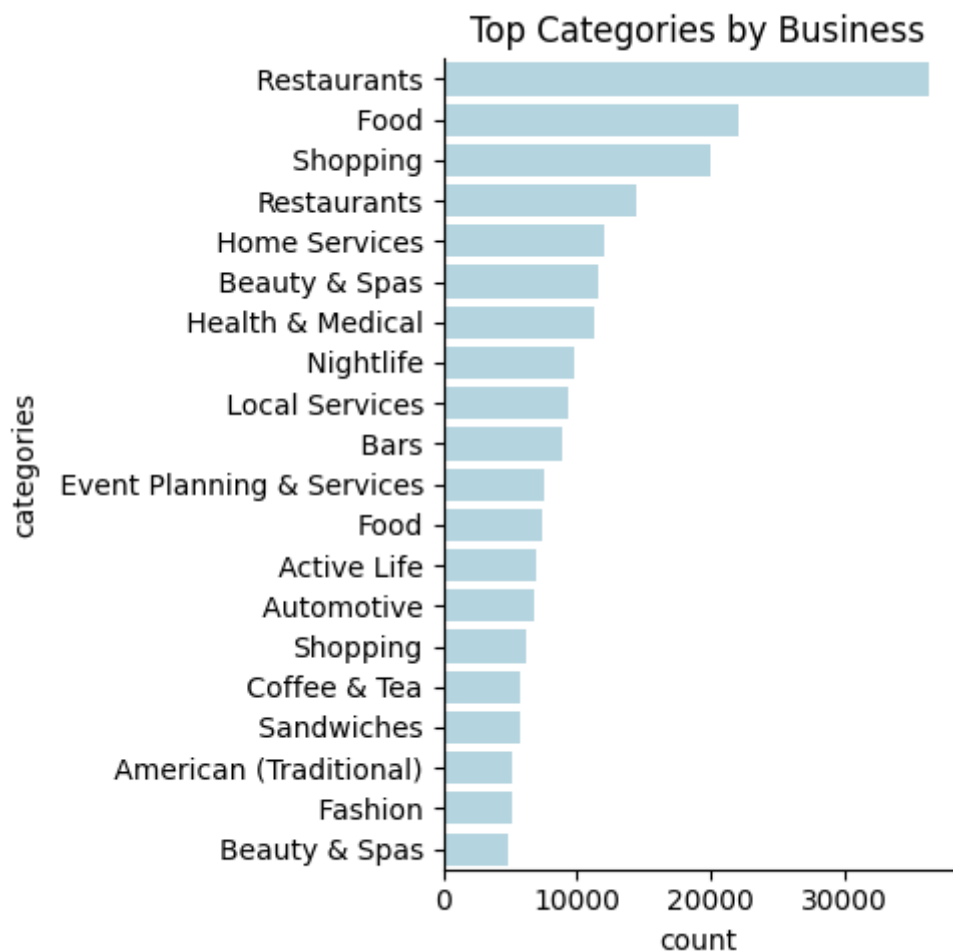
```
In [14]: business_group = business_cat_explode.groupby('categories')
business_cat_bar = business_group.count().orderBy('count',ascending=False)
business_cat_bar.show(20)
```

categories	count
Restaurants	36340
Food	22094
Shopping	20056
Restaurants	14423
Home Services	12001
Beauty & Spas	11633
Health & Medical	11390
Nightlife	9808
Local Services	9299
Bars	8914
Event Planning &...	7617
Food	7375
Active Life	7039
Automotive	6785
Shopping	6149
Coffee & Tea	5735
Sandwiches	5697
American (Tradit...	5235
Fashion	5231
Beauty & Spas	4941

only showing top 20 rows

```
In [15]: bcdf = business_cat_bar.toPandas()
```

```
In [16]: sns.catplot(y='categories',x='count', data=bcdf[0:20], kind='bar',color='lightbl
plt.title('Top Categories by Business')
plt.tight_layout()
plt.show()
%matplotlib plt
```



Yelp User Reviews Analysis

```
In [17]: review.printSchema()
```

```
root
|-- business_id: string (nullable = true)
|-- cool: long (nullable = true)
|-- date: string (nullable = true)
|-- funny: long (nullable = true)
|-- review_id: string (nullable = true)
|-- stars: double (nullable = true)
|-- text: string (nullable = true)
|-- useful: long (nullable = true)
|-- user_id: string (nullable = true)
```

```
In [18]: review.select('business_id', 'stars').show(5)
```

```
+-----+-----+
| business_id | stars |
+-----+-----+
| buF9druCkbuXLX526... | 4.0 |
| RA4V8pr014UyUbDvI... | 4.0 |
| _sS2LBIGNT5NQb6PD... | 5.0 |
```



```
|0AzLzHfOJgL7R0whd...| 2.0|
|8zehGz9jnxPqXtOc7...| 4.0|
+-----+-----+
only showing top 5 rows
```

In [19]:

```
#calculate user review on business grouping by business

rev_avg = review.select('business_id', 'stars') \
    .withColumn('stars', col('stars').cast('float')) \
    .groupBy('business_id') \
    .agg(mean('stars').alias('avg_stars'))
```

In [20]:

```
rev_avg.show(5)
```

```
+-----+-----+
|          business_id|          avg_stars|
+-----+-----+
|OXcBg_6vgi-J3nZzI...| 3.852272727272727|
|bxy3khT-2R66tcdKj...| 3.986344537815126|
|f3teByaeIKPTYetAa...| 4.329608938547486|
|3KqpiLDAjeeMmZeU-...| 2.555555555555554|
|_hkHeU2cBH9fzthId...| 3.168195718654434|
+-----+-----+
only showing top 5 rows
```

In [21]:

```
#joining business and rev_avg df to compare average user rating to business rating

bus_avg_star = business.join(rev_avg, business.business_id == rev_avg.business_id)
```

In [22]:

```
bus_avg_star.select('avg_stars', 'stars', 'name', 'city', 'state').show(5)
```

```
+-----+-----+-----+-----+-----+
|          avg_stars|stars|          name|          city|state|
+-----+-----+-----+-----+-----+
|1.9090909090909092| 2.0|          Safeway|      Vancouver|WA|
|2.966292134831461| 3.0|Cracker Barrel Ol...|Pickerington|OH|
|2.81981981981982| 3.0|Peaceful Restaurant|      Vancouver|BC|
|                5.0| 5.0|      ATX Architects|          Austin|TX|
|4.524271844660194| 4.5|    Evergreen Eatery|          Boston|MA|
+-----+-----+-----+-----+-----+
only showing top 5 rows
```

Calculating Skew Between Business Rating and Average User Rating on said Business

To calculate skewness between user rating and business rating, I will calculate % difference.

(Average Rating - Business Rating) / Business Rating


Above formula will find the percent change between average user rating and business rating

```
In [23]: skew = bus_avg_star.withColumn('skew', (col('avg_stars')-col('stars')) / col('st
```

```
In [24]: skewdf = skew.toPandas()
```

```
In [25]: plt.clf()
plt.figure(figsize=(12, 6))
sns.distplot(skewdf['skew'])
plt.title('User Review Skew Relative to Business Rating')
plt.tight_layout()

plt.show()
%matplotlib plt
```



The above histogram shows that the average user rating is typically lower than the business rating due to left skew.

Rating Comparison of Elite vs Normal Users on Same Business

```
In [26]: #joining review and user df to classify which review is considered elite members
review_user = review.join(user, review.user_id == user.user_id)
```

```
In [27]: review_user.printSchema()
```

```
root
|-- business_id: string (nullable = true)
```

```

-- cool: long (nullable = true)
-- date: string (nullable = true)
-- funny: long (nullable = true)
-- review_id: string (nullable = true)
-- stars: double (nullable = true)
-- text: string (nullable = true)
-- useful: long (nullable = true)
-- user_id: string (nullable = true)
-- average_stars: double (nullable = true)
-- compliment_cool: long (nullable = true)
-- compliment_cute: long (nullable = true)
-- compliment_funny: long (nullable = true)
-- compliment_hot: long (nullable = true)
-- compliment_list: long (nullable = true)
-- compliment_more: long (nullable = true)
-- compliment_note: long (nullable = true)
-- compliment_photos: long (nullable = true)
-- compliment_plain: long (nullable = true)
-- compliment_profile: long (nullable = true)
-- compliment_writer: long (nullable = true)
-- cool: long (nullable = true)
-- elite: string (nullable = true)
-- fans: long (nullable = true)
-- friends: string (nullable = true)
-- funny: long (nullable = true)
-- name: string (nullable = true)
-- review_count: long (nullable = true)
-- useful: long (nullable = true)
-- user_id: string (nullable = true)
-- yelping_since: string (nullable = true)

```

In [28]:

```
#calculation of elite user average rating on business
```

```

norm_avg = review_user.filter(review_user.elite=='').withColumn('stars', col('stars')
    .groupBy('business_id') \
    .agg(mean('stars').alias('norm_avg')))

```

In [29]:

```
#calculation of normal user average rating on business
```

```

elite_avg = review_user.filter(review_user.elite.contains(',')').withColumn('stars',
    .groupBy('business_id') \
    .agg(mean('stars').alias('elite_avg')))

```

In [30]:

```
#joining both df from above to get ready to calculate skewness
```

```
avgjoined = norm_avg.join(elite_avg, norm_avg.business_id == elite_avg.business_id)
```

In [31]:

```
avgjoined.show(3)
```

```

+-----+-----+-----+-----+
+

```

business_id	norm_avg	business_id	elite_avg
-36nnCT71XE0InJXK...	1.903225806451613	-36nnCT71XE0InJXK...	2.0
-QO103c2B22yi_On0...	2.6470588235294117	-QO103c2B22yi_On0...	3.4
-VVUUPK0ytYjpJ_S7...	2.814814814814815	-VVUUPK0ytYjpJ_S7...	2.8333333333333335

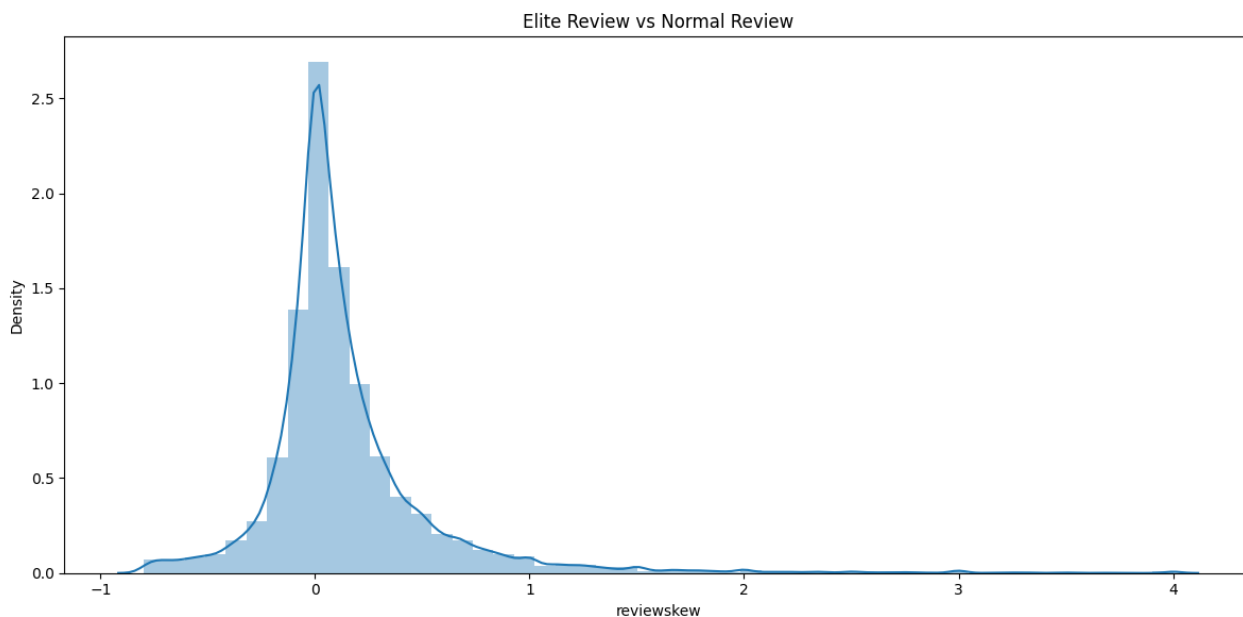
only showing top 3 rows

```
In [32]: #calculating % difference between elite and normal reviews
reviewskew = avgjoined.withColumn('reviewskew', (col('elite_avg')-col('norm_avg')
```

```
In [33]: reviewskewdf = reviewskew.toPandas()
```

```
In [34]: plt.clf()
plt.figure(figsize=(12, 6))
sns.distplot(reviewskewdf['reviewskew'])
plt.title('Elite Review vs Normal Review')
plt.tight_layout()

plt.show()
%matplotlib plt
```



Above histogram shows that Elite members typically rate business higher than normal users due to right skew.